

AI-Generated Content Transparency Compliance Requirements (Watermarking)

EU AI Act, Art. 50(2)

1. Executive Summary

Cornerstone must implement AI-generated content transparency mechanisms for **all AI systems (“AIS”)**, including general-purpose AI systems, that **generate synthetic audio, image, video, or text content**.

The EU AI Act (“**AIA**”) requires Cornerstone to ensure that AI outputs (“**Output**”) generated or manipulated by AIS are marked in a machine-readable format and detectable as artificially generated or manipulated, as far as technically feasible (all measures together also referred to as “**Watermarking**”). For Cornerstone, this applies to AI features that generate or materially manipulate text, image, audio, or video Outputs for customers.

Compliance deadline: August 2, 2026 (new products, not only new feature to an existing product) / December 2, 2026 (existing products)

2. Definition of AI System

The AIA defines AIS as *a machine-based system that is designed to operate with varying levels of autonomy and that may exhibit adaptiveness after deployment, and that, for explicit or implicit objectives, infers, from the input it receives, how to generate outputs such as predictions, content, recommendations, or decisions that can influence physical or virtual environments.*

3. Requirements:

- **Mark AI-generated or AI-manipulated Outputs:** For every AI feature that generates synthetic audio, image, video, or text content, the **Output must be marked in a machine-readable format so that it can be detected as artificially generated or manipulated.**
- **Ensure technical robustness:** The **marking solution must be effective, interoperable, robust, and reliable** as far as technically feasible, taking into account the type of content, technical limitations, implementation costs, and the generally acknowledged state of the art.
- **Apply to relevant generative AI and GPAI-powered features:** Teams must assess all Generative AI, AI agent, Assistive AI, and Machine Learning features to determine

whether they generate or substantially manipulate text, image, audio, or video Outputs. The requirement applies regardless of whether the feature is built directly by Cornerstone or relies on an AIS / AI Model integrated into a Cornerstone product.

- **Preserve existing user-facing AI notices where applicable:** Human-readable notices remain appropriate where Output is shown to users, but Art. 50(2) additionally requires technical marking / detectability of AI-generated or manipulated Outputs.
 - For generative AI chat or text Outputs, continue using the agreed footer text below where applicable, however that **does not suffice for compliance** with Art. 50(2) AIA: **“Generated by AI. Check for accuracy”**.
- **Multi-layer approach:** Teams shall implement a layered approach where technically feasible, combining **complementary measures such as machine-readable metadata, watermarking/embedding in the Output, and appropriate detection or verification mechanisms**. The exact layers may differ by content type and product scenario, but the selected approach must be capable of demonstrating that the Output is detectable as AI-generated or AI-manipulated.
- **Document exceptions carefully:** The marking obligation does not apply where the AI system only performs an assistive function for standard editing or does not substantially alter the input data or its semantics. If a team believes an exception applies, Compliance must be involved and **approve** on a case-by-case basis before the feature is treated as out of scope.
- **Notifying** the user about the AI generation is not sufficient to comply with Art. 50(2) AIA. The provision aims to protecting the “information ecosystem”, i.e. it aims at making the text (machine) detectable as AI generated when it is disseminated to third parties.
- **Additional Resources:** This [study](#) published by the EU provides further technical details or can be of help. Further guidance can be found [here](#) and [here](#).
- **“Free form text”:** **For free form text (not contained in a file) it is sufficient to only implement one layer of Watermarking for detection.**
- **Length:** The watermarking requirement does not apply for very short texts of less than **200 tokens**.
- **Human modification of Output:** The Watermarking requirement does not need to be maintained where the final Output is no longer substantially AI-generated or AI-manipulated because a human has materially changed the Output. **As a practical threshold, the marking should only be dropped where the human contribution has substantially altered the input data or its semantics, i.e. the final content reflects meaningful human authorship or editorial responsibility rather than merely light**

edits, formatting changes, shortening, translation, or stylistic adjustments. If the AI-generated part remains material to the final Output, or if there is uncertainty whether the human changes are sufficient, the Watermarking should be retained and the exception must be documented and approved by Compliance.

- **Model Providers:** Please note, model providers have **not** made public any reliable concepts of Watermarking used in their models yet and therefore we cannot rely on them.
- **Technical feasibility:** The Guidance emphasizes that a marking solution must be technically feasible. This means we are not obliged to make use of a technical solution that is not yet developed or available on the market, or that is technically unfeasible for implementation. Also, there is no requirement to implement the latest research still in an experimental stage. **However, this area moves fast, and our customers won't accept us telling them we cannot do it!**

4. Evidence / Submission Requirement:

- A standardized approach for the various scenarios (text documents, videos, free-form text, etc.) that equally applies to and is standardized across all products must first be agreed with Legal/AI Governance.
- The agreed approach then needs to be implemented **for all** features that provide Output as described above. Exceptions need to be flagged to mboudreau@csod.com and need to be approved and documented by AI Governance.

5. Questions:

Please reach out to Marina Boudreau with any questions: mboudreau@csod.com